

EQUITY RESEARCH — INITIATION OF COVERAGE

NVIDIA CORPORATION

NASDAQ: NVDA

The Growth Is Real. The Question Is What Discount Rate You Apply to It.

Rating **HOLD**

Price target (blended) \$179 (12-18 month view)

Price as of August 28, 2026 close (\$217.43)

Prepared for Private Equity / Equity Research portfolio

EXHIBIT 0.1 Snapshot

\$217.43 SHARE PRICE	\$5.24T MARKET CAP	\$179 BLENDED PT	-17.6% PT vs. PRICE
\$96.2bn Q2 FY27 REVENUE	+106% Q2 REVENUE YoY	75.0% Q2 GROSS MARGIN	27.5x TRAILING P/E

CONTENTS

- 01** Recommendation, Up Front
- 02** Company Overview
- 03** Financial Performance
- 04** Investment Thesis — The Bull Case
- 05** Key Risks — The Bear Case
- 06** Valuation — Discounted Cash Flow
- 07** Trading Comparables
- 08** Catalysts — What to Watch
- 09** Final Verdict & Scorecard
- 10** Key Limitations of This Analysis
- 11** Sources

01 Recommendation, Up Front

THE ONE-LINE CALL

HOLD. Near-term growth is about as close to a sure thing as this business gets — the real disagreement, and the entire reason this stock is hard to value, is what discount rate an investor should demand for a company this large, this fast-growing, and this dependent on a handful of customers' capital-spending decisions.

NVIDIA is not a hard company to underwrite on growth. Revenue for the quarter ended July 26, 2026 was \$96.2 billion, up 106% year-over-year, with Data Center revenue of \$89.0 billion, up 117%. The company has beaten its own revenue guidance for thirteen consecutive quarters. Management has already guided Q3 FY2027 to \$108 billion and characterized full fiscal-year 2028 as growing roughly 70% — a figure it explicitly frames as supply-constrained, with underlying demand running closer to 100%. None of that is in serious analytical dispute.

What is genuinely in dispute is the price. A standalone discounted cash flow using NVIDIA's own reported beta (2.23, reflecting real historical volatility including two sharp drawdowns) implies a fair value of roughly \$133 per share — 39% below where the stock trades. The same cash flows, discounted at a lower rate that treats NVIDIA's near-term growth as substantially de-risked by its disclosed \$279 billion of supply and capacity commitments, imply roughly \$226 — modestly above the current price. The stock trades almost exactly at the point where the bull case for the discount rate becomes the only assumption doing the work. That is not a reason to be bearish, and it is not a reason to be bullish — it is the reason this is a HOLD rather than a conviction call in either direction.

My rating: HOLD, blended price target ~\$179 (the average of the two DCF scenarios, which I weight far more heavily than the P/E-based comp cross-check — see Section 7 for why that comp is unreliable in NVIDIA's specific case right now).

02 Company Overview

NVIDIA designs GPUs and the surrounding networking, software, and systems stack that together power accelerated computing — originally for graphics rendering, and increasingly for the training and inference of AI models at data-center scale. The business now reports through two platforms: Data Center (compute and networking for hyperscalers, AI-native companies, enterprises, and sovereign customers) and Edge Computing (the legacy gaming/graphics business plus automotive and robotics), following a 2026 presentation change that recast prior-period comparisons.

EXHIBIT 2.1 Q2 FY2027 revenue by platform

Platform	Q2 FY27 Revenue	QoQ	YoY
Data Center	\$89.02bn	+18%	+117%
Hyperscale	\$48.71bn	+13%	+102%

Platform	Q2 FY27 Revenue	QoQ	YoY
AI Clouds, Industrial & Enterprise	\$40.31bn	+25%	+138%
Edge Computing	\$7.20bn	+13%	+27%
Total revenue	\$96.22bn	+18%	+106%

Source: NVIDIA Q2 FY2027 Form 10-Q, revenue by market platform table, quarter ended July 26, 2026.

Data Center is now 92.5% of total revenue, up from a business that was still meaningfully diversified across gaming and professional visualization as recently as three years ago. That concentration is the central fact governing both the bull and bear cases in this report: it means NVIDIA's fortunes are now tied overwhelmingly to hyperscaler and AI-native capital spending decisions, and to the durability of its architectural lead (Blackwell Ultra currently, Vera Rubin ramping through Q3 FY2027, expected to represent roughly 20% of Data Center revenue in that quarter) over both GPU competitors and hyperscalers' growing internal custom-silicon programs.

The company's moat is usually described in three layers: raw chip performance, the NVLink/InfiniBand networking stack that lets thousands of GPUs act as one system, and CUDA — the software layer that has become the default programming model for GPU-accelerated computing over nearly two decades, creating real switching costs for developers and enterprises alike. All three layers matter more than any single one; a competitor matching NVIDIA on chip benchmarks alone (which several already claim to, in narrow workloads) still faces the networking and software gap.

03 Financial Performance

EXHIBIT 3.1 Three-year financial summary, fiscal years 2024–2026

\$bn except per-share	FY2024A	FY2025A	FY2026A
Total revenue	\$60.9	\$130.5	\$215.9
Revenue growth	n/a	114.2%	65.4%
GAAP gross margin	72.9%	75.0%	71.1%
Operating income (approx.)	\$33	\$81	\$129
Operating margin	54.1%	62.4%	59.7%
Net income (approx.)	\$30	\$73	\$106
Diluted EPS	\$1.19	\$2.94	\$4.90

Source: NVIDIA Form 10-K filings, FY2024–FY2026 (fiscal year ends the last Sunday of January). FY2026 gross margin declined from FY2025's peak on rising memory input costs — the same dynamic management now guides will pressure margins further into Q4 FY2027 before pricing actions restore them in FY2028.

The FY2026 gross-margin dip (75.0% to 71.1%) is worth sitting with rather than skating past, because it previews the single clearest near-term risk in this report: NVIDIA doesn't fully control its input costs, and memory (HBM) pricing has been rising industry-wide as every AI infrastructure buyer competes for the same constrained supply. Management's own Q2 FY2027 guidance makes the trajectory explicit — gross margin bottoming around 71–72% in Q4 FY2027 before recovering to 72–73% in FY2028 as NVIDIA's own price increases take effect. That's a company

telling investors, in its own words, that a cost pressure is currently outrunning its ability to pass costs through, and that the fix is a price increase still ahead of it, not behind it.

Balance sheet position has changed in a way many investors haven't fully priced into their mental model of this company: NVIDIA moved from a large net-cash position in FY2024–FY2026 to net debt of roughly \$11 billion by Q2 FY2027 (cash and short-term investments of \$22.4bn against \$33.4bn of total debt), reflecting debt-funded capacity investment to keep pace with demand. That shift is immaterial to the valuation given NVIDIA's \$5.2 trillion market cap, but it's a genuine change in capital structure worth noting for anyone still working from an older mental model of NVIDIA as a debt-free cash machine.

04 Investment Thesis — The Bull Case

4.1 Demand visibility is unusually high for a company this size

The strongest bull argument isn't really about AI as a category — it's about the specificity of NVIDIA's near-term revenue. Unlike most high-growth technology companies, where next year's revenue is a forecast built on trend extrapolation and market-sizing assumptions, NVIDIA's near-term revenue is substantially a function of already-disclosed supply and capacity commitments: approximately \$279 billion, mostly memory secured for the Vera Rubin ramp. A company that has already locked in the inputs for its next several quarters of production is, by construction, carrying less demand-side uncertainty than its historical volatility (and beta) would suggest — which is exactly the argument for why a bull-case investor should apply a materially lower discount rate than NVIDIA's trailing five-year beta implies.

4.2 The AWS deal and the broadening of the customer base

An expanded compute agreement with AWS for 2 million additional GPUs, announced August 27, 2026, added roughly the value of a mega-cap company to NVIDIA's market cap in a single session (+8.74%). The magnitude of that single-day reaction is itself informative about how the market is currently treating incremental hyperscaler commitments — as close to pure upside, with little embedded skepticism. That's a useful signal of sentiment, not necessarily of fundamental value, and it cuts both ways: it confirms hyperscaler demand remains unsatisfied, but it also confirms how little room for disappointment is priced into the stock at current levels.

4.3 Optionality beyond the current AI-training narrative

Robotics, autonomous vehicles, and “sovereign AI” (national governments building independent AI infrastructure) are all real, disclosed growth vectors that this report's explicit DCF forecast does not meaningfully underwrite — the base and bull cases both are effectively pure extrapolations of the current Data Center training/inference business. To the extent any of these newer categories scale meaningfully within the explicit forecast window, both DCF scenarios in this report are, if anything, conservative rather than aggressive on the revenue side.

05 Key Risks — The Bear Case

5.1 Customer concentration and the capex-cycle question

NVIDIA's growth is now a function of perhaps a dozen companies' capital-spending decisions — the major hyperscalers and a handful of AI-native labs. Q2 FY2027's combined hyperscaler capex ran at roughly \$166.0 billion, up 87% year-over-year; that pace of increase cannot continue indefinitely, if only on the arithmetic of what share of revenue and free cash flow the largest technology companies can sustainably redirect to infrastructure. The bear case does not require AI to be a bubble in the pejorative sense — it only requires hyperscaler capex growth to normalize from 87% toward something more sustainable, which would mechanically slow NVIDIA's growth well before this report's FY2029–FY2031 explicit forecast assumes it does.

5.2 Competition is real, even if not yet revenue-relevant

Broadcom's AI-related revenue reflects its position as the primary beneficiary of hyperscalers' growing appetite for custom silicon (ASICs) designed to reduce dependence on a single GPU supplier charging 75% gross margins. The bear case for NVIDIA isn't that any competitor matches its current chip performance — none clearly do — it's that at NVIDIA's margin structure, the economic incentive for every large customer to fund an alternative is enormous and growing, and custom-silicon share of total AI compute has reportedly been expanding faster than prior projections.

5.3 Geopolitical and export-control exposure

Data Center shipments of Hopper products to China were less than 1% of Data Center revenue in Q2 FY2027, and NVIDIA's own guidance now assumes zero China compute revenue going forward — a policy-driven headwind that has already cost the company a \$4.5 billion charge in an earlier quarter (Q1 FY2026, tied to H20 export licensing) and remains a binary, headline-driven risk that no amount of financial modeling fully captures.

5.4 Margin and multiple compression, together

The most underappreciated bear case isn't growth deceleration in isolation — it's growth deceleration arriving at the same time gross margin is already under pressure from memory costs, which would compress both the numerator (earnings growth) and the multiple the market is willing to pay for that growth simultaneously. That combination, rather than either factor alone, is historically how richly-valued growth stocks experience their sharpest drawdowns.

06 Valuation — Discounted Cash Flow

6.1 Why two scenarios, not one

A single-point DCF would misrepresent the actual state of the debate on this stock. The disagreement among genuinely informed investors is not primarily about NVIDIA's near-term revenue (which is substantially guided) or its long-run competitive position (which is a real but separate debate) — it is about what discount rate is appropriate for a company whose realized historical volatility (beta 2.23) reflects a smaller, more speculative version of itself, but whose current revenue is backed by disclosed supply commitments most speculative growth stocks never have. This report builds both cases explicitly rather than picking one and hiding the other.

EXHIBIT 6.1 *Revenue build — base case vs. bull case, FY2027E–FY2031E*

\$bn	FY2027E	FY2028E	FY2029E	FY2030E	FY2031E
Revenue growth — base case	87.5%	70.0%	35.0%	20.0%	13.0%
Revenue growth — bull case	87.5%	70.0%	45.0%	30.0%	20.0%
Revenue — base case	\$404.8	\$688.2	\$929.1	\$1,114.9	\$1,259.9
Revenue — bull case	\$404.8	\$688.2	\$998.0	\$1,297.4	\$1,556.9

FY2027E built from Q1+Q2 FY27 actuals (\$81.6bn + \$96.2bn), the Q3 FY27 guide (\$108.0bn), and an analyst estimate for Q4 FY27 (~\$119bn). FY2028E uses management's own “~70%, supply-constrained” full-year guidance directly. FY2029E–FY2031E are this analysis's own moderation — no company guidance exists that far out, and holding 70–100% growth indefinitely at this revenue base is not mathematically sustainable against any plausible total addressable market.

6.2 The WACC debate, made explicit

EXHIBIT 6.2 WACC derivation — base case vs. bull case

Component	Base Case	Bull Case
Risk-free rate (10Y UST)	4.1%	4.1%
Equity risk premium	5.5%	5.5%
Beta	2.23 (as reported)	1.60 (analyst assumption)
Cost of equity / WACC	16.37%	12.90%

Base-case beta is NVIDIA's actual reported 5-year beta (Google Finance, Aug 2026). Bull-case beta of 1.60 is this report's own assumption, not a reported figure — it reflects the argument that realized historical volatility overstates NVDA's forward risk given its current revenue visibility, while still sitting well above 1.0 given genuine customer-concentration and export-control risk that a pure infrastructure-utility beta would not carry.

6.3 DCF output

EXHIBIT 6.3 Valuation bridge — both scenarios

Component	Base Case	Bull Case
Enterprise value	~\$3,164bn	~\$5,383bn
(+) Net cash / (net debt)	(\$11bn)	(\$11bn)
Equity value	~\$3,153bn	~\$5,372bn
Diluted shares (bn)	24.10	24.10
Implied value per share	\$132.95	\$225.51
Current price	\$217.43	\$217.43
Implied premium / (discount)	(38.9%)	+3.7%

\$132.95	\$225.51	\$179.23
-----------------	-----------------	-----------------

BASE CASE DCF	BULL CASE DCF	BLENDED (AVG.)
---------------	---------------	----------------

Read together, the two scenarios bracket the current price almost symmetrically — the stock trades close to, but modestly below, what the bull case justifies, and well above what the base case justifies. That is a textbook description of a stock priced for its bull case to be substantially correct, with limited room for the base case's more conventional risk-pricing to reassert itself without a meaningful drawdown. Full formula detail, including a two-way WACC/terminal-growth sensitivity table for each scenario, is in the accompanying Excel model.

07 Trading Comparables

EXHIBIT 7.1 Semiconductor / AI infrastructure peer multiples

Company	Beta	P/E (TTM)	P/E (Fwd)	EV/EBITDA
NVIDIA (NVDA)	2.23	27.48x	n/a	n/a
AMD	2.49	120.79x	42.61x	79.87x
Broadcom (AVGO)	1.47	61.83x	23.55x	43.08x

AMD and Broadcom statistics per stockanalysis.com (S&P Global Market Intelligence), data as of Aug 24–28, 2026.

The single most counter-intuitive fact in this entire report sits in that table: NVIDIA, the largest and most dominant company in the group, trades at a lower trailing P/E (27.5x) than either AMD (120.8x) or Broadcom (61.8x). That is not evidence NVIDIA is cheap — it's evidence that AMD and Broadcom's trailing earnings are currently depressed relative to their own growth expectations, which inflates their trailing multiples mechanically, while NVIDIA has already grown into a large share of the earnings the market is paying for. Applying the peer-average trailing P/E (91.3x, dominated by AMD's distorted figure) to NVIDIA's own TTM EPS produces an implied price above \$400 — a number I'm not using anywhere in this report's actual valuation conclusion, because it's an artifact of a noisy peer average, not a reliable read on NVIDIA's own fair value. This is the same category of mistake as blindly averaging in a merger-distorted P/E in a comp set, and it's worth naming explicitly rather than quietly excluding without explanation.

The more useful comps takeaway is qualitative: Broadcom's much lower beta (1.47) despite also being a major AI-infrastructure beneficiary suggests the market already discounts Broadcom's earnings stream as more diversified and lower-risk than NVIDIA's — consistent with Broadcom's business spanning networking, software, and custom silicon across many customers, versus NVIDIA's concentration in merchant GPUs sold overwhelmingly to a small hyperscaler customer set.

08 Catalysts — What to Watch

- Q3 FY2027 earnings (expected late November 2026) — tests the \$108bn guide and, more importantly, whether gross margin holds near the guided 74.0% or shows further memory-cost pressure.

- Vera Rubin ramp through Q3–Q4 FY2027 — the architecture is guided to reach ~20% of Data Center revenue in Q3; a slower or faster ramp than that directly moves the FY2028 growth number this entire valuation leans on.
- Any further hyperscaler capex guidance updates — the single most direct read-through to NVIDIA's own growth, given how concentrated its customer base is.
- China policy developments — current guidance assumes zero China compute revenue; any change in either direction is a real, non-trivial swing factor not modeled as upside in this report's base or bull case.
- Custom-silicon disclosures from major hyperscalers — the clearest leading indicator of the bear case's competitive-share thesis actually materializing in the numbers, as opposed to remaining a commentary-only risk.

09 Final Verdict & Scorecard

EXHIBIT 9.1 Coverage scorecard — five criteria, scored independently

Criterion	Score (/10)	Basis
Near-term growth visibility	9	Q3 FY27 substantially guided; FY28 guided directly by management at ~70%, framed as supply- rather than demand-constrained
Competitive moat durability	8	CUDA + networking stack remains a genuine, multi-year switching-cost advantage; custom silicon is a real but not yet revenue-relevant threat
Valuation vs. base-case DCF	4	Current price sits 39% above the base-case (reported-beta) DCF — requires the bull-case discount-rate argument to be correct
Customer concentration risk	5	A dozen or so hyperscaler/AI-native customers now drive the overwhelming majority of revenue — structurally higher risk than a diversified customer base
Margin trajectory	6	Gross margin currently under pressure from memory costs, guided to bottom in Q4 FY27 before recovering in FY28 — credible, but not yet proven out

FINAL VERDICT

HOLD, blended price target ~\$179 (12–18 month view). NVIDIA is one of the best-run, highest-moat businesses in technology, and its near-term growth is about as well-supported as hyper-growth ever gets. But the stock's current price only clears a base-case, textbook-CAPM valuation if you accept the bull case's central assumption — that NVIDIA's forward risk is now meaningfully lower than its own trading history suggests. That's a defensible view, not an obviously correct one, and a portfolio built on it should size the position accordingly rather than treat the growth story alone as sufficient justification for the price.

Would I initiate a full position here? No — not because the business is anything other than excellent, but because the valuation leaves very little margin of safety if the bull case's discount-rate argument turns out to be wrong, or if hyperscaler capex growth normalizes faster than this report's already-aggressive FY2029–FY2031 assumptions allow for. The more interesting entry point, in my view, is any pullback toward the \$150–180 range that isn't accompanied by a genuine deterioration in the Q3/Q4 FY2027 guided numbers — that would be a market re-rating the discount rate without the underlying growth story actually having changed, which is exactly the setup a disciplined investor should want to buy into rather than chase at the top of the range.

10 Key Limitations of This Analysis

- FY2029–FY2031 revenue growth in both DCF scenarios is this report's own judgment, not company guidance — no public guidance exists that far out, and the true long-run growth/deceleration path for a company at this scale and growth rate is inherently uncertain.
 - The bull-case beta of 1.60 is an analyst assumption, explicitly flagged as such, not a reported or consensus figure — reasonable investors could argue for a beta anywhere between the reported 2.23 and this report's 1.60 bull-case assumption.
 - FY2026 operating income and net income in Section 3 are approximated from reported EPS and share count rather than pulled line-by-line from the 10-K's income statement; directionally accurate, not to the last decimal.
 - The peer comp set (AMD, Broadcom) is narrow; a fuller comp set (e.g., TSMC, Marvell, custom-silicon-exposed names) might tell a richer story about how the market is pricing AI-infrastructure exposure more broadly, and was not built out here given the P/E-comp reliability issues already discussed in Section 7.
 - This report does not attempt a sum-of-the-parts valuation separating Data Center from Edge Computing, even though the two now have very different growth and margin profiles — a natural next step if this project extends further.
-

11 Sources

-
- NVIDIA Form 10-K, fiscal years 2024–2026 (SEC EDGAR)
 - NVIDIA Form 10-Q and Form 8-K press release, Q2 fiscal year 2027 (SEC EDGAR, Aug 26, 2026)
 - NVIDIA Q4 & Fiscal 2026 earnings press release (Feb 25, 2026) and Q1 FY2027 outlook statement (May 20, 2026)
 - Google Finance, NVDA key statistics (price, beta, shares outstanding, P/E) — Aug 28, 2026
 - stockanalysis.com / S&P Global Market Intelligence — AMD and Broadcom statistics, Aug 24–28, 2026
 - S&P Global Market Intelligence, “Nvidia earnings preview: Q2 2027”
 - Various financial media coverage of the Aug 26–28, 2026 earnings release and AWS compute agreement, cross-checked across multiple sources for consistency